

Jetson Orin Nano Latency — native TensorRT, FP32 vs FP16 (batch size 1, 96×96)

