

Jetson Orin Nano — native TensorRT latency (batch=1, 96×96, MAXN_SUPER, all 16 models)

