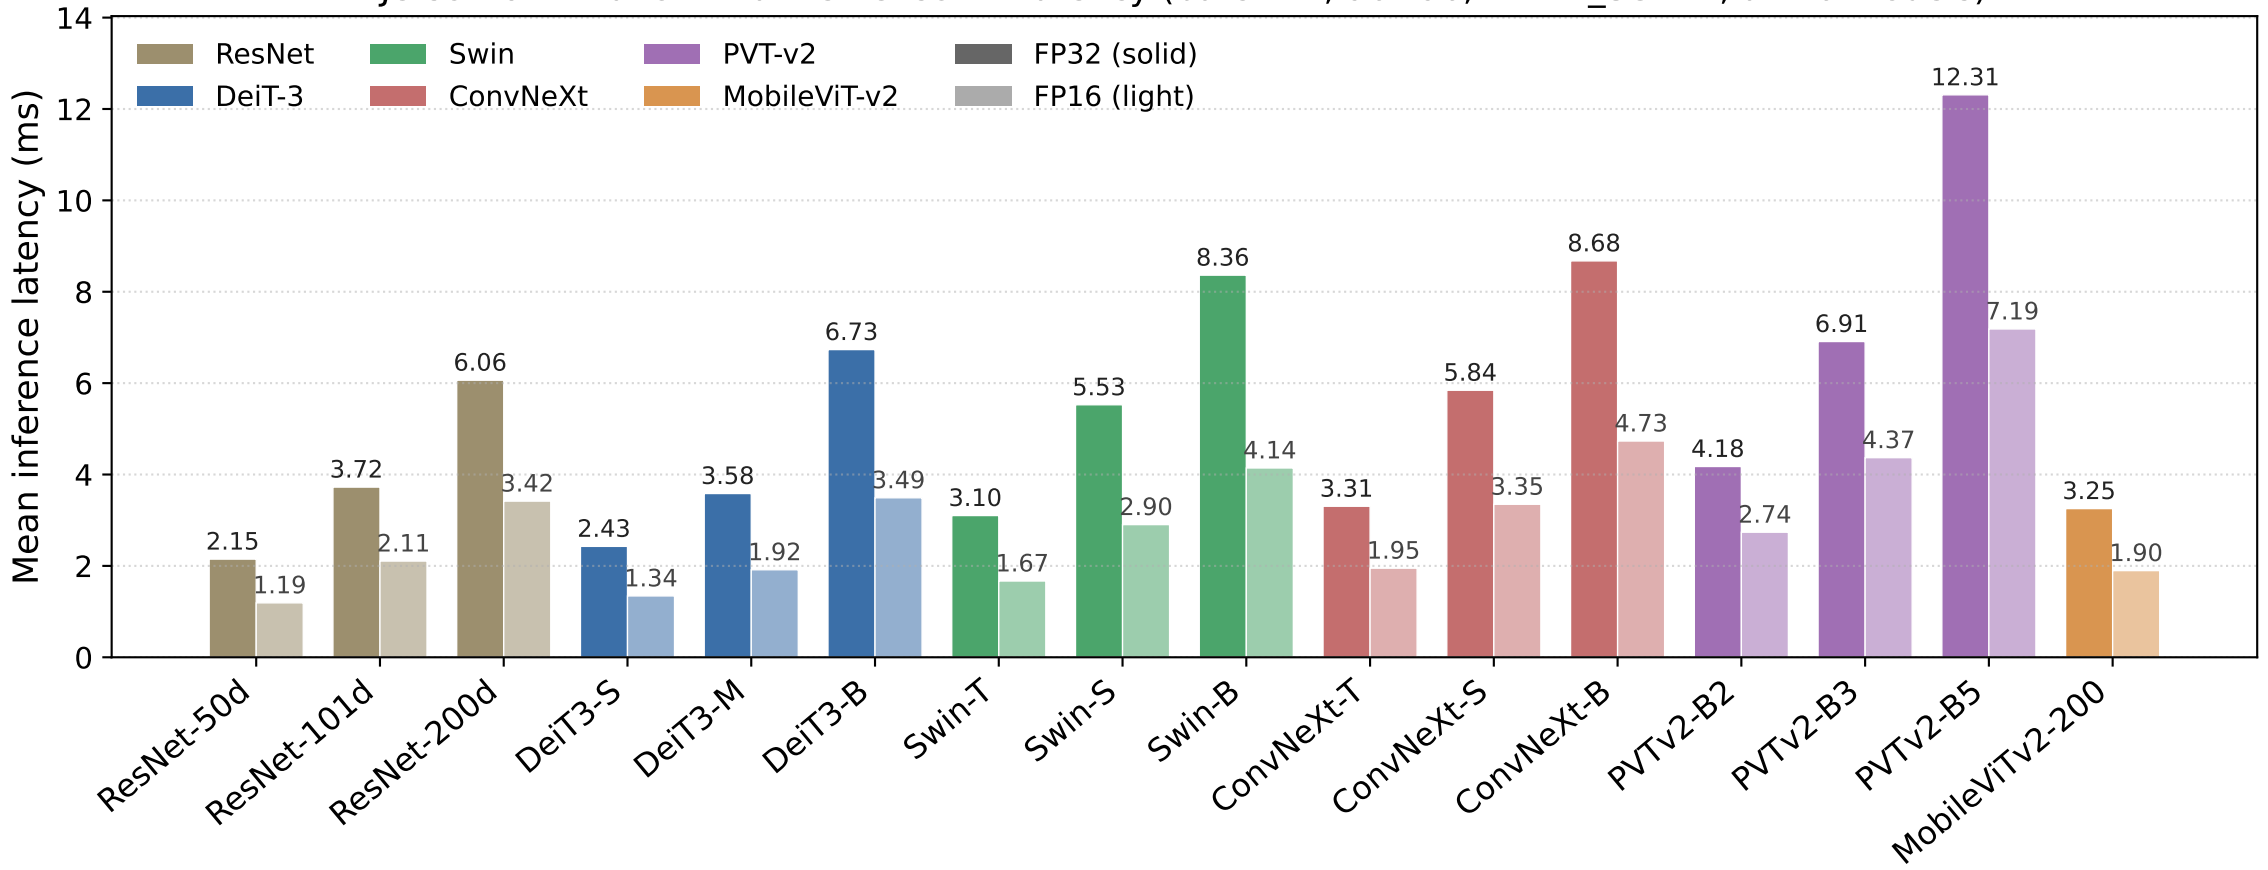
Jetson Orin Nano — native TensorRT latency (batch=1, 96×96, MAXN_SUPER, all 16 models)
Mean inference latency (ms)
0
2
4
6
8
10
12
14
ResNet
DeiT-3
Swin
ConvNeXt
PVT-v2
MobileViT-v2
FP32 (solid)
FP16 (light)
2.15
1.19
3.72
2.11
6.06
3.42
2.43
1.34
3.58
1.92
6.73
3.49
3.10
1.67
5.53
2.90
8.36
4.14
3.31
1.95
5.84
3.35
8.68
4.73
4.18
2.74
6.91
4.37
12.31
7.19
3.25
1.90
ResNet-50d
ResNet-101d
ResNet-200d
DeiT3-S
DeiT3-M
DeiT3-B
Swin-T
Swin-S
Swin-B
ConvNeXt-T
ConvNeXt-S
ConvNeXt-B
PVTv2-B2
PVTv2-B3
PVTv2-B5
MobileViTv2-200